

About PulsePoint:

[PulsePoint](#) is a fast-growing AdTech company using real-time data to transform healthcare. We help brands and agencies interpret the hard-to-read signals across their customers' health journey to produce the most complete view of the customer. Our award-winning platforms use machine learning to activate this data, enabling easy, instantaneous marketing, predictive analytics, and decision support.

In this role, you will:

- Design, develop, and maintain LLM-based AI components to improve the efficiency of internal teams and user-facing communications, spanning from data retrieval to chat interfaces.
- Experiment and evaluate new frameworks, models, and architectures emerging in the area.
- Improve the underlying data infrastructure to address the challenges posed by LLM usage.
- Partner with cross-functional teams of analysts, software engineers, and user-facing departments to identify relevant use cases and design appropriate solutions.
- Contribute to the technical vision for the foundation of generative AI systems in PulsePoint.

You will be a part of the Core Engineering team, which is responsible for the development and maintenance of the PulsePoint Exchange and Life Advertising platform:

- Handles hundreds of thousands of transactions per second, billions of times each month.
- Evaluates, selects, and optimizes ad-serving based on advanced statistics and machine learning.
- Returns responses collected from dozens of parties in milliseconds.
- Constantly evolves to meet market demands that change in days and weeks, not months/years.
- Factors thousands of data points in every serving decision.

We're small enough that you can own something and have a direct impact, but big enough that you don't have to do it alone. We are flexible enough to accommodate your interests as they evolve and find the best application for your passion at the moment.

The company and the team have a flat hierarchy, with technical management writing code and direct access to business, product, and operations.

Lastly, our engineers have lots of empowerment and freedom of action (but we don't water down our responsibilities or expectations).

Basic Qualifications:

- Python experience, experience with basic data manipulation packages like numpy and pandas, solid SQL skills, and understanding TDD principles.
- Independence, initiative, commitment, and responsibility to what you have tamed.
- Excellent problem-solving and critical thinking: you will need to figure out your own mistakes, and sometimes not just yours, especially in such a new field.

Preferred Qualifications:

- Experience developing, delivering, and supporting AI and ML algorithms or technologies (e.g. LLM Inference, Embeddings, and VectorDBs, ranking algorithms, and time series analysis).
- Experience with k8s, as we are mostly self-hosted.
- Java & Python proficiency.